


Break-Even Volatility

Nicolas Mitoulis

A dissertation submitted to the Faculty of Commerce, University of Cape Town, in partial fulfilment of the requirements for the degree of Master of Philosophy.

February 10, 2019

*MPhil in Mathematical Finance,
University of Cape Town.*



Declaration

I declare that this dissertation is my own, unaided work. It is being submitted for the Degree of Master of Philosophy at the University of Cape Town. It has not been submitted before for any degree or examination at any other University.

February 10, 2019

Abstract

A profit or loss ($P\&L$) of a dynamically hedged option depends on the implied volatility used to price the option and implement the hedges. Break-even volatility is a method of solving for the volatility which yields no profit or loss based on replicating the hedging procedure of an option on a historical share price time series. This dissertation investigates the traditional break-even volatility method on simulated data, how the break-even formula is derived and details the implementation with reference to MATLAB. We extend the methodology to the Heston model by changing the reference model in the hedging process. Resultantly, the need to employ characteristic function pricing methods arises to calculate the Heston model sensitivities. The break-even volatility solution is then found by means of an optimisation of the continuously delta hedged $P\&L$ over the Heston model parameters.

Acknowledgements

I would like to express my gratitude to my supervisors at the African Institute of Financial Markets and Risk Management, Professor David Taylor and Obeid Mahomed, for their guidance throughout the MPhil degree. I would also like to thank Emlyn Flint for proposing the topic of this dissertation. Lastly, my appreciation is extended to my family and friends for all the support they have given me.

Contents

1. Introduction	1
2. Break-Even Volatility	3
2.1 Basic methodology via discrete delta hedging	3
2.2 Traditional BEV formula	5
2.2.1 Derivation	5
2.2.2 BEV as a Γ -weighted average	7
2.3 Implementation notes	8
2.3.1 Fixed-point iteration	8
2.3.2 Uniqueness of solution	8
2.3.3 Aggregation of results	9
2.3.4 Caveats	10
2.4 Shortcomings	11
3. Stochastic Volatility	13
3.1 The Heston stochastic volatility model	13
3.2 Monte Carlo experiment with standard BEV	15
3.3 Proposed method using Heston Greeks	17
4. Implementation and Results	20
4.1 Standard BEV	21
4.2 Optimising over Heston parameters	23
4.3 Incorporating realised volatility	27
5. Discussion and Conclusion	29
Bibliography	31
A. Model specifics	32
A.1 Black-Scholes PDE	32
A.2 The Heston stochastic volatility model	32
A.2.1 Simulation	32
A.2.2 Characteristic function pricing	33
A.2.3 The Heston Greeks	34

List of Figures

2.1	Three-month GBM price path with delta hedged $P\&L$ of a short call ($K = 100$) as a function of σ and corresponding break-even volatility skew.	4
2.2	A GBM price path (with dynamics as in Section 2.1), a call option's $P\&L$ as function of σ and volatility skew plot.	11
3.1	True implied volatility (transparent) and break-even volatility estimated surface (10 000 paths, MSE: 3.7873e-07).	17
4.1	Optimising the whole surface: True implied volatility surface (transparent) and break-even volatility surface, averaged over all options (50 paths). MSE = 3.0074e-05.	21
4.2	Point-wise optimisation: True implied volatility surface (transparent) and break-even volatility surface, zeroing each path's $P\&L$ for a specific option separately (50 paths), in between 3 standard deviation bounds (red). MSE = 2.2235e-05.	22
4.3	Break-even volatility distributions for two options of 1-year term (sample size: 50).	23
4.4	Point-wise optimisation: True implied volatility surface (transparent) and point-wise optimised for Heston parameters, zeroing each options' average $P\&L$ (50 paths). MSE = 2.6976e-05.	25
4.5	MSE as a function of sample size for the 3-month volatility skew, with strike price range 85% – 115% of S_0	25
4.6	Optimising the whole surface: True implied volatility surface (transparent) and full surface optimised for Heston parameters, averaging all options' $P\&L$'s (50 paths). MSE = 8.6556e-06.	27
4.7	True implied volatility surface (transparent) and point-wise optimised over Heston parameters (50 paths) using the realised variance in place of the actual variance path. MSE = 1.0397e-04.	28

List of Tables

4.1	Break-even volatility estimates for two options (50 sample paths). . .	23
4.2	Parameter estimates for 3-month skew with increasing sample size. .	26

Chapter 1

Introduction

Traditional option volatility estimation involves calculating the annualised standard deviation of daily realised log returns from a historical time series of stock price data. This is only sensible in light of the constant volatility assumption of geometric Brownian motion (GBM). Naively, the realised volatility over a chosen period of time can be used in the Black-Scholes formula to yield a price, for any given strike or maturity. This ignores any market skew, where the implied volatility (the volatility input which equates the Black-Scholes option price to the market price) is not constant for options with different strikes and maturities on the same underlying.

A dynamically hedged option recoups the difference between realised volatility and the implied volatility, this is captured as either a profit or a loss. To do this, one must implement the correct hedge and an accurate method of estimating the option volatility. The traditional estimate of volatility gives no indication of how the market is actually pricing options on the underlying as it is not contingent on the option's strike or maturity. This makes the differential between the implied and realised volatility difficult to discern.

Dupire (2006) presents a method using dynamic (delta) hedging to calculate break-even volatility (BEV) surfaces. He explains that this estimate is the fair volatility to both sides of an option contract, as it is the volatility which zeroes the option's profit or loss. The paper highlights BEV surfaces as a tool to understand market volatility surfaces. Dupire assesses the level of evident mispricing and severity of market skews by comparing the implied volatilities to what the fair break-even volatility would have been over the option's life. Subsequently, break-even volatility has evolved from a historical analysis tool to become a volatility estimation procedure, mostly for longer maturity options or where similar option information is unavailable.

It is difficult to obtain a realised volatility estimate for different strikes or maturities. This is precisely what the break-even volatility methodology aims to achieve.

It is shown in Section 2.2 from the calculation of BEV, any profit or loss arising during a dynamically hedged option is influenced by the option's Gamma and is thus linked to strike and time to expiry. Solely based on past data, the methodology creates option specific, constant volatility estimates representing how the market should have been pricing and hedging if it were fair. Comparing these estimates with implied volatilities gives an indication of market sentiment and the degree of mispricing on an option.

The standard BEV methodology makes no assumptions about the data generating process, it merely examines a time series of historical stock price data to produce an estimate of an option's volatility. The method entails retrospectively delta hedging an option over a historical stock price path using the Black-Scholes model and solving for the volatility which makes the hedge a fair game. However, if we have additional information about which type of process is generating the data, changing the reference model to match should in theory result in an improved hedge. With improved replication of the profit/loss of the option, we would then find a better break-even estimate of the volatility used to price and hedge the option.

The stochastic volatility model of Heston (1993) introduces an interesting extension of the break-even volatility methodology. The standard method of break-even volatility is data intensive and typically employs statistical bootstrapping of historical stock returns to simulate more price time series and increase the data size. This cannot be done with a stochastic volatility process where the time series must stay unbroken to maintain the serial correlation structure. The Heston model is commonly used and makes for a challenging adaptation to the standard methodology due to its high parameter dimensionality and associated data problems for the break-even methodology.

Break-even volatilities are constant estimates for each option, much like implied volatilities, and do not provide a term-structure of volatilities. In essence, they serve as a base pricing platform. It is therefore important for an accurate initial price/volatility to be estimated, by employing an appropriate reference model, not necessarily Black-Scholes. This dissertation reviews and examines the standard BEV methodology, providing notes of its implementation in MATLAB and highlighting its shortcomings. In an attempt to improve BEV estimates, we analyse the effect of changing the reference model from Black-Scholes to Heston. This is done by a comparison of each method's efficacy at returning the true implied volatility surface of the Heston process which generated the data.

Chapter 2

Break-Even Volatility

The former financial market model of choice has been the Black-Scholes model to calculate an initial option premium and subsequent hedges, where historical volatility has been used to determine the volatility σ . Not only does this assume the future will be as volatile as the past, the constant volatility assumption necessitates that options on the same underlying will be priced with the same volatility. Due to many market models being used and in most cases when the assumption of constant volatility is violated, the historical volatility estimate would produce a premium unsuitable to the volatility risk and hedging cost.

The attempt to find an accurate price or volatility for an option leads to a more critical problem of finding what the fair volatility to both sides of the option contract should be, this would be the volatility which causes the option to break even. [Suzuki and Vyas \(2011\)](#) note that in empirical tests (where the stock price process is known) using the historical volatility over the life of an option or even the market implied volatility at the time were insufficient in matching the total hedging costs without market friction. The report focuses on extending the break-even methodology to incorporate transaction costs and other market conventions, as well as showing break-even volatility to be an efficient and fair pricing platform.

2.1 Basic methodology via discrete delta hedging

In his seminal paper on break-even volatility, [Dupire \(2006\)](#) formalised the methodology as a means of finding the fair market volatility skew or surface. The volatility estimates which form the skew are solved for and equate the profit or loss ($P\&L$) of a dynamically hedged option to zero.

This can be visualised in a toy example using a simulated GBM price path with dynamics

$$dS_t = 0.08S_t dt + 0.2S_t dW_t$$

Suppose the current price of a share is 100, the risk-free rate r is 0.06 and we've sold

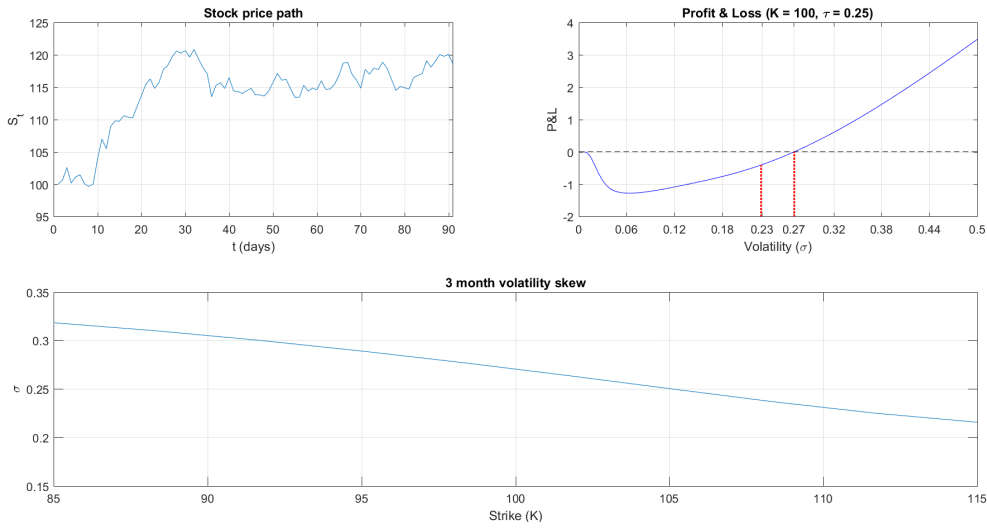


Fig. 2.1: Three-month GBM price path with delta hedged $P\&L$ of a short call ($K = 100$) as a function of σ and corresponding break-even volatility skew.

a call option on the share with strike $K = 100$ and a term of three months. After performing a daily delta hedging experiment on this data until maturity using the Black-Scholes model and a chosen constant volatility, the hedge profit or loss then depends on the choice of volatility used to price and hedge the option. The solution which zeroes the hedge is the break-even volatility.

In Figure 2.1 the option expires deep in the money as seen in the top left three-month price path. By calculating the $P\&L$ as a function of the chosen volatility σ , we construct the graph on the top right of Figure 2.1. From the $P\&L$ graph, the break-even volatility is 0.27, which is considerably higher than the volatility used to simulate the data (0.2). The break-even estimate seems sensible due to the option expiring deep in the money, as such a volatile stock path requires the initial premium to be inflated as well as to calculate the necessary delta holdings in the stock. Even the historical volatility over the period of 0.23 is not sufficient to allow the hedge to break-even. Regardless of the level of moneyness, this historical estimate would have been the same and is thus a misleading estimate with which to price and hedge due to its inability to account for option specifics (strike and term). This lends weight to the use of break-even volatility estimates which in essence is a reweighting of the asset returns used in the historical estimate, where the weights depend on strike and maturity.

By finding the break-even volatility for different moneyness levels, one obtains a volatility skew all from this three-month price path. This skew is naturally highly dependent on the path used to create it and one could obtain an unrealistic view

of how an option should be priced by simply observing one price path. In Figure 2.1 the stock price increases substantially within three months, lower strikes correspond to the option being deeper in the money and thus a higher premium must be charged to account for this as depicted in the skew.

This process can be repeated for different option maturities to create an entire volatility surface for a given asset. Results can then be smoothed by averaging over a number of historical periods. This can be done as follows:

1. Specify a strike K and maturity τ , delta hedge (using Black-Scholes) an option daily over a period in history and compute the $P\&L(\sigma)$.
2. Find the break-even volatility $\sigma_{K,\tau}$ such that $P\&L(\sigma_{K,\tau}) = 0$.
3. Repeat 1 and 2 over all strikes and terms to create the volatility surface for that specific period.
4. Repeat above steps over all historical periods and average the results to arrive at the final break-even volatility surface.

By implementing the process above, a separate surface is created for each historical period which is useful to see prices or volatilities changing over time. This means one obtains a distribution of volatility surfaces from the individual distributions of break-even volatilities for each option. The break-even volatility surface is usually found by averaging the surfaces, alternatively for a single option the average break-even volatility would be the final estimate.

2.2 Traditional BEV formula

Discrete delta hedging a historical price path can be used to find break-even volatilities, however this method is not often used by practitioners. In this section, we show a method of deriving a formula for the $P\&L$ which stems from a continuous delta hedging argument. In practice a discretised version of this formula is used to find break-even volatilities, where an integral is approximated with a summation.

2.2.1 Derivation

Suppose the true dynamics of an underlying stock price process are given by

$$dS_t = \mu_t S_t dt + \sigma_t S_t dW_t$$

where μ_t and σ_t are non-negative adapted/predictable processes. Since this is a general process that we assume for the stock and no assumptions are put on the

form of μ_t and σ_t , break-even volatility is seen as a non-parametric volatility estimation procedure.

Assuming a constant risk-free rate r , we sell and hedge a European option on the underlying using the Black-Scholes model with constant volatility σ .

Let $Y_t = C(t, S_t)$ be the Black-Scholes price process of this option. Let X_t be the value of the hedge portfolio $\Theta = (\theta_t^{(0)}, \theta_t^{(1)})$, consisting of the bank account growing at the risk-free rate and the stock, where $\theta_t^{(1)} = \frac{\partial C}{\partial S}(t, S_t) = \Delta_t$.

By accumulating the hedging error $Z_t := X_t - Y_t$ to maturity T , we derive an approximate formula for the delta-hedged option's profit/loss. To do this, we first need the dynamics of the hedge portfolio and of the Black-Scholes price.

Using

$$\begin{aligned} X_t &= \theta_t^{(0)} e^{rt} + \Delta_t S_t \\ \implies \theta_t^{(0)} &= e^{-rt} (X_t - \Delta_t S_t) \end{aligned}$$

to ensure the portfolio is self-financing, the SDE for the hedge portfolio is

$$\begin{aligned} dX_t &= \theta_t^{(0)} r e^{rt} dt + \theta_t^{(1)} dS_t \\ &= r(X_t - \Delta_t S_t) dt + \Delta_t (\mu_t S_t dt + \sigma_t S_t dW_t) \\ &= (rX_t + \Delta_t (\mu_t - r) S_t) dt + \Delta_t \sigma_t S_t dW_t \end{aligned}$$

Using Itô's lemma on the Black-Scholes price process and substituting the option price sensitivities for the standard Greek letters gives us

$$\begin{aligned} dY_t &= dC(t, S_t) = \frac{\partial C}{\partial t}(t, S_t) dt + \frac{\partial C}{\partial S}(t, S_t) dS_t + \frac{1}{2} \frac{\partial^2 C}{\partial S^2}(t, S_t) d[S]_t \\ &= \Theta_t dt + \Delta_t (\mu_t S_t dt + \sigma_t S_t dW_t) + \frac{1}{2} \Gamma_t \sigma_t^2 S_t^2 dt \\ &= (\Theta_t + \Delta_t \mu_t S_t + \frac{1}{2} \Gamma_t \sigma_t^2 S_t^2) dt + \Delta_t \sigma_t S_t dW_t \end{aligned}$$

Now $C(t, S_t)$ satisfies the Black-Scholes PDE (shown in Appendix A.1) with constant volatility σ and risk-free rate r . The price sensitivities in the first term above are dictated by the pricing model (Black-Scholes in this case) and by substituting for Θ_t from the Black-Scholes PDE allows dY_t to be written as

$$dY_t = (rY_t + \Delta_t (\mu_t - r) S_t + \frac{1}{2} \Gamma_t (\sigma_t^2 - \sigma^2) S_t^2) dt + \Delta_t \sigma_t S_t dW_t$$

The expression for dZ_t is then found by eliminating the like terms in dX_t and dY_t

$$\begin{aligned} dZ_t &= dX_t - dY_t = r(X_t - Y_t) dt - \frac{1}{2} \Gamma_t (\sigma_t^2 - \sigma^2) S_t^2 dt \\ &= rZ_t dt + \frac{1}{2} \Gamma_t S_t^2 (\sigma^2 - \sigma_t^2) dt \end{aligned}$$

Solving with e^{-rt} as an integrating factor and integrating from 0– T , the $P\&L$ at maturity is obtained as

$$\begin{aligned} e^{-rt}dZ_t - re^{-rt}Z_tdt &= \frac{1}{2}e^{-rt}\Gamma_t S_t^2(\sigma^2 - \sigma_t^2)dt \\ \int_0^T d(e^{-rt}Z_t) &= \int_0^T \frac{1}{2}e^{-rt}\Gamma_t S_t^2(\sigma^2 - \sigma_t^2)dt \\ Z_T &= \frac{1}{2} \int_0^T e^{r(T-t)}\Gamma_t S_t^2(\sigma^2 - \sigma_t^2)dt \end{aligned} \quad (1)$$

In the final step shown above, $Z_0 = 0$ due to the hedge portfolio being initially equal to the Black-Scholes price. Naturally the formula for the $P\&L$ in Equation (1) is approximated by a daily sum using market stock price data. We can acquire the instantaneous variance σ_t^2 with the relationship $(\frac{dS_t}{S_t})^2 = \sigma_t^2 dt$ and obtain a formula for the $P\&L$ in terms of the volatility σ used to price and hedge as

$$P\&L(\sigma) = \frac{1}{2} \sum_{i=0}^{N-1} e^{r(T-t_i)}\Gamma_{t_i} S_{t_i}^2 \left(\sigma^2 \Delta t_i - \left(\frac{\Delta S_{t_i}}{S_{t_i}} \right)^2 \right) \quad (2)$$

where $t_N = T$ is the maturity of the option, $\Delta t_i = \delta t$ are all equal daily increments and $\Delta S_{t_i} = S_{t_{i+1}} - S_{t_i}$.

A strength of the break-even methodology is it is non-parametric in nature, it makes no assumptions as to the true underlying stock price process. This is outright stated in the derivation of Equation (1), where we assume general stock dynamics and drift and volatility processes. By using the Black-Scholes model to price and hedge the option does not mean it is assumed the stock follows GBM, the Black-Scholes model acts as a converter in the same way the model is used to transform a price into an implied volatility.

To construct a volatility surface, one follows the same procedure as in the discrete hedging case and finds the volatility which zeroes Equation (2) for a specific option. This formula yields similar, if not slightly more stable, results than the discrete hedging method. Additionally, when concisely implemented, it is not as computationally expensive.

2.2.2 BEV as a Γ -weighted average

If we equate Equation (2) to zero and solve for σ we arrive at the following equation

$$\sigma^2 \delta t = \sum_{i=0}^{N-1} w_i \left(\frac{\Delta S_{t_i}}{S_{t_i}} \right)^2 \quad \text{where } w_i = \frac{e^{r(T-t_i)}\Gamma_{t_i} S_{t_i}^2}{\sum_{j=0}^{N-1} e^{r(T-t_j)}\Gamma_{t_j} S_{t_j}^2} \quad (3)$$

We see that the break-even volatility is a weighted average of squared returns or realised variances. For a specific volatility i.e specific strike and term, the aver-

age will be dominated by the returns where the Γ_t values are highest. This happens for price changes close to the strike price and towards maturity of the option; these price changes have arguably the greatest effect on the outcome of the option. Thus, points on a break-even skew are averages of the same returns with different weights due to Γ 's dependence on the strike. The downside of this is that an average of break-even volatilities for a specific option term leads back to the historical volatility for that period. Constructing the skew is then susceptible to sample variation because it depends on the chosen period. A different period is seen as a different sample and will result in a different estimate of the break-even volatility skew in the same way historical volatility changes for each period. This is to be considered when obtaining an estimate, however it is of little consequence as the method's aim is to form an estimate of where the level of the market skew should have been.

2.3 Implementation notes

2.3.1 Fixed-point iteration

In Equation (3), the σ for which we are trying to solve is needed to calculate the Γ_{t_i} values. Thus if one chooses to find break-even volatility using this equation, then a fixed-point iterative algorithm (e.g. Newton's method) is needed.

Likewise for Equation (2), an iterative algorithm is best for accuracy. In the interest of speed and if outright accuracy is not required (i.e. if the break-even volatility is a base estimate which will be adjusted), a simple grid search can be used to find the intercept region where the $P\&L$ is zero and interpolated to the desired accuracy.

2.3.2 Uniqueness of solution

In the implementation of the standard methodology, for complete accuracy, we use a combination of a grid search and MATLAB's `fzero` function which employs the bisection method (amongst other iterative methods) to find the root of Equation (2). Dupire (2006) mentions that the root may not be unique, but a strictly positive root is guaranteed. We see this in Figure 2.1 where zero is also a solution. In some cases it appears that the $P\&L$ is non-negative, such as in Figure 2.2. This is due to machine precision where extremely small negative values of the $P\&L$ function are returned as zero. A fine enough grid search is also not feasible as this becomes too computationally expensive when repeated for multiple options. To combat this, the MATLAB optimizer `fmincon` is used alongside the grid search and root finding

algorithms, specifically for these occurrences. When a root cannot be found, the solution is then the largest σ such that the $P\&L$ is zero.

2.3.3 Aggregation of results

Time windows

Once a period in history is chosen for which the break-even volatility levels will be calculated, the next step is to subdivide this time series into smaller time windows with length corresponding to the term of the option. This presents another possibility as to whether these time windows are overlapping or not. The BEV methodology performs greater and smoother results are obtained with a larger dataset, and overlapping time windows might allow this. Consequently, the break-even volatilities for each time window will no longer be identically and independently distribution (*i.i.d.*) as they will be averages of the same squared returns. This would render the distribution of break-even volatilities misleading and may not be appropriate for finding interval bound estimates for the break-even volatility of an option.

In practice, it is common to use independent time series to find each break-even volatility estimate and assemble a distribution or average to obtain a final estimate. In doing this, one must either rebase the stock price at the beginning of each time series to a common value or alternatively work with relative option strike prices.

Another commonly used method of data aggregation involves both statistical bootstrapping and Monte Carlo simulation. One can convert a time series of share prices into a time series of returns

$$S = \{S_0, S_1, S_2, \dots, S_N\} \implies R = \{r_1, r_2, r_3, \dots, r_N\}$$

where $r_i = \log(\frac{S_i}{S_{i-1}})$ for $i \in \{1, 2, \dots, N\}$. With sufficient evidence of *i.i.d.* returns, one can randomly sample from R and create a new time series of returns $R^* = \{r_1^*, r_2^*, \dots, r_M^*\}$ where each $r_i^* \in R$ and M corresponds to the option term under consideration i.e. $t_M = T$ the option maturity. The time series of returns R^* is used to create a new sample path of share prices

$$S^* = \{S_0^*, S_1^*, S_2^*, \dots, S_N^*\} \quad \text{where} \quad S_0^* = S_0 \quad \text{and} \quad S_i^* = S_{i-1}^* \exp(r_i^*)$$

This can be repeated to create multiple sample paths, each with a break-even volatility. This is similar to bootstrapping where the initial sample size of the returns is increased and then Monte Carlo simulation is performed to the desired accuracy by repeatedly sampling returns to create share price paths to find break-even volatilities and the results averaged. Likewise, with the subdivision into independent time series, it is important to base the initial value of each generated path to a common value as we have done with S_0 above.

Averaging methods

When the break-even volatility is found for each independent period of history, a distribution of volatilities is created for a specific option from the accumulation of these estimates and the final estimate is the mean of the distribution.

Another method is to average the option's $P\&L$'s for all historical periods considered and find the single volatility, which applies to all periods, that zeroes this average. In tests using simulated data, this method provides benefits in computational efficiency, yields smoother surfaces and converges faster to the true implied volatility surface. The drawback of this method is that we will no longer have a distribution of break-even volatilities for each point on a skew/surface to provide a certain level of confidence in an estimate.

In finding a break-even volatility estimate, we would like this volatility to imply an option price such that the option $P\&L$ is zero on average. In simply averaging periodical break-even volatilities instead, we're zeroing each $P\&L$ but with different volatilities all for the same option. Thus it makes sense to apply one volatility to all historical periods so that the average $P\&L$ is nil.

A reason why the latter method may be more desirable, we are forcing the distribution of the $P\&L$'s to be centred around zero for a given volatility. It seems that an outlying share price time series would have an associated break-even volatility which skews the volatility distribution more so than the corresponding $P\&L$ can influence the break-even volatility estimate which zeroes the average $P\&L$. This is especially evident in short maturities where a few large returns may dominate the weighted average in Equation (3) in finding the volatility, thus causing a radical estimate. Adding to this, large negative and positive $P\&L$'s from different time periods may offset each other linearly in the average whereas these $P\&L$'s may not translate linearly to their respective break-even volatilities so that the average volatility is equivalently balanced.

2.3.4 Caveats

As stated previously, machine precision may play a part in making the root finding more difficult and tedious, thus one must make use of a combination of algorithms to account for this. For instance, we noticed in a few cases the $P\&L$ function repeatedly returns slightly positive values before going back to zero, for increasing σ values. This is precisely what arises in the $P\&L$ plot in Figure 2.2, causing initial root finding to be incorrect. This may not occur often and is dependent on the volatility of the historical period, but is likely to occur when using many sample historical periods or bootstrap sampling more paths. This can cause inconsisten-

cies in the break-even volatility surfaces, thus setting up the algorithm and stress testing with simulated data goes a long way.

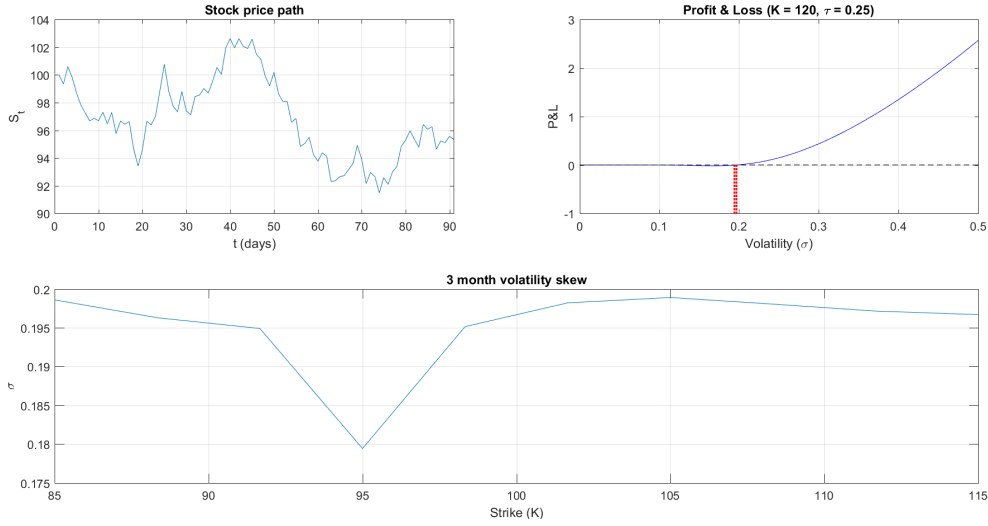


Fig. 2.2: A GBM price path (with dynamics as in Section 2.1), a call option's $P\&L$ as function of σ and volatility skew plot.

There may be instability in break-even estimates for options with short maturities, for strikes well outside of the price range of a given path and even for strikes which are well within or at-the-money (ATM) as seen in the volatility skew in Figure 2.2 for a strike of 95. Fortunately averaging over many time windows helps smooth this out but it contributes to slow convergence of shorter maturity options to the true implied volatility skew.

The next point is not necessarily a caveat, it is more as something to note. Nothing has been said of the issue of the risk-free rate. The correct daily rate along the time series of prices may be input each day in the hedging process, however we found that changes in the constant risk-free rate as high as 10% did not heavily impact break-even estimates and the volatility skews marginally change. This is in line with the findings of Dupire (2006) and he notes that the traditional historical estimate also disregards interest rates.

2.4 Shortcomings

The break-even methodology is predominantly a pricing platform for instruments which trade on volatility and exotic or long-term options. It is also a means of gauging the volatility premium currently offered by the market. As Dupire (2006) notes, the aim is not to produce a term structure of volatilities which would aid in

hedging, but rather an initial price.

The slow convergence and long run-times make this methodology difficult to work with and set up. One is forced to use a long time series of prices in order to have enough data to overcome instabilities in the break-even estimate. However, if a large dataset is used, these break-even volatility surfaces become static through time. If the historical period stretches too far into the past, the estimates would not be susceptible to frequent change as more data emerges and this could make it a troublesome pricing technique in a continually changing market.

A finite sample may be taken and bootstrapped in order to increase the size of the dataset and ensure smoother results. A problem with this, which we will see in the Heston model, is that by bootstrapping returns enforces these new prices to be log-normally distributed. This Monte Carlo style approach causes the volatility surface to flatten as in a Black-Scholes world ([Dupire, 2006](#)) and converge to the realised volatility of the sample.

Chapter 3

Stochastic Volatility

Stochastic volatility (SV) models have been widely adopted in practice for their ability to produce similar volatility skews observed in the market. This is most notable for equity options which have negatively sloping volatility skews caused by share price returns being negatively correlated with volatility. This is known as the leverage effect and has been empirically observed by [Bakshi *et al.* \(1997\)](#), amongst others.

SV models shed the restrictive constant volatility assumption in the Black-Scholes world and allow the volatility to vary through time as a stochastic process. Volatility has also been observed to cluster, i.e. periods of large price changes yielding high realised volatility as well as periods of low volatility. The choice of model used in these periods dictates the pricing and hedging errors, and a model which assumes constant volatility would exacerbate these errors. To reduce this, the volatility assumption must be continually changed to react to the market, however this causes disorganised changes in the hedge ratios ([Gatheral, 2011](#), p. 2). The BEV methodology involves a constant volatility assumption for all the hedges, therefore it is reasonable to update the calculation with another choice of model.

3.1 The Heston stochastic volatility model

In this dissertation we examine the model of [Heston \(1993\)](#) in our extension of the standard BEV methodology. The Heston model allows the stock price to be dependent on a time-varying variance process and produces skews and return distributions consistent with those seen in the market.

Model dynamics

For a stock price S_t and variance process ν_t at time t , the dynamics of the Heston model are given by the following stochastic differential equations:

$$\begin{aligned} dS_t &= \mu S_t dt + \sqrt{\nu_t} S_t dW_t^{(1)} \\ d\nu_t &= \kappa(\theta - \nu_t)dt + \sigma \sqrt{\nu_t} dW_t^{(2)} \end{aligned}$$

where $W_t^{(1)}, W_t^{(2)}$ are correlated Brownian motions.

The parameters of the Heston model are given by:

- μ Expected/real-world rate of return.
- κ Mean reversion rate.
- θ Long-term variance level.
- σ Volatility of the variance process.
- ρ Correlation parameter where $dW_t^{(1)}dW_t^{(2)} = \rho dt$.

As in the Black-Scholes case, μ is the drift of the stock which is replaced by r in the risk-neutral world. The parameter κ controls the rate of mean reversion of the variance to the long-term level of ν_t , given by θ . The magnitude of κ would determine the length of time the volatility spends above or below θ . This is consistent with the common observation of volatility tending to cluster for certain periods in the market. Intuitively, we must have $\kappa, \theta > 0$ for the process to be mean-reverting. These conditions ensure that the variance will revert to the long-term level, because the drift term for $d\nu_t$ will be positive for periods of low volatility and likewise the variance will be dragged down when ν_t is above θ as there will be negative drift. The parameter ρ allows for correlation between the stock and variance processes, and σ determines the volatility of the latter process. These parameters control the height and skewness of the return distribution. With a negative correlation parameter, the return distribution is skewed to the left. The fatter left tail means lower returns are more likely than higher returns, thus the option price for increasing strikes will decrease. This creates the negative volatility skew and explains the leverage effect.

The variance process is known as the square root process, introduced by [Cox et al. \(1985\)](#), and is strictly positive if $2\kappa\theta > \sigma^2$ (known as the Feller condition), given the initial volatility is nonnegative i.e. it cannot become negative. Since the variance is a latent process where the level cannot be seen in the market, the initial variance ν_0 becomes another parameter of the model.

Reasons for considering the Heston model for Break-even Volatility

There has not been any substantial research into the performance of the break-even volatility methodology on stochastic volatility data. The Heston model provides an assessment of the BEV methodology's ability to recover the theoretically correct volatility skew.

The need to bootstrap sample share price returns to simulate new price time series is vital to the methodology in order to construct smooth and stable surfaces. This method is commonly used by practitioners, but can no longer be done with Heston simulated data or any market data which may follow a SV process. The correlation between price and volatility creates a serial correlation structure amongst the daily prices and thus the day-to-day returns will also share this correlation. The inherent assumption that the daily returns are *i.i.d.* involved in bootstrapping the data and sampling daily returns to create new paths would destroy this correlation and lead to a flat volatility surface. The inability to do this in practice would leave one with limited data when looking for a recent market volatility skew. Alternatively, one could choose to use a large time series of prices with the drawback that the volatility surface will not be susceptible to change as more data emerges.

Dupire (2006) mentions that alternative models to Black-Scholes may be used in order to improve the replication/hedge and then these "volatilities", which are parameters from the alternative model, may be converted into Black-Scholes implied volatilities. We examine how the standard methodology performs on a simulated dataset without conducting bootstrap sampling to create more data. We compare this to using the Heston model option price sensitivities (Γ in Equation (2)), which we hope will provide an improvement in the form of reduced estimation error and will somewhat overcome the sampling error involved with a small dataset.

3.2 Monte Carlo experiment with standard BEV

As a sanity check, we perform a Monte Carlo simulation experiment. This is to verify that, given enough data, the standard BEV methodology can recover the true volatility surface implied by a Heston model with specific parameters. This being possible, the mean squared error (MSE) would decrease by using the theoretically correct model in the hedging process. This is not unequivocally checked due to the amount of data in this simulation being computationally infeasible for calculating the $P\&L$ with the Heston model.

The following parameter set was used in this simulation experiment which will become the base scenario for the rest of this dissertation:

- $\nu_0 = 0.06$
- $\kappa = 9$
- $\theta = 0.06$
- $\sigma = 0.5$
- $\rho = -0.4$

We set $S_0 = 100$ and hedge with risk-free rate $r = 0.06$. We let $\mu = 0.08$ for the sake of simulating non-risk neutral price paths. This, however, makes minimal difference to the BEV estimates and we are interested solely in the listed parameters for pricing purposes where μ has no impact.

Here we are performing a pure simulation experiment, therefore there is no need to consider subdivision of time series or bootstrapping daily returns, each path will be a complete and independent Heston price path simulated via a Milstein scheme (see Appendix A.2.1 for this discretisation). We simulate 10 000 daily stock price paths of the longest maturity considered (1.5 years), where BEV estimates for shorter-dated options are based on the initial segment, with length corresponding to the specific option, of each independent path. The $P\&L$'s are calculated using Equation (2) and the method of averaging all $P\&L$'s for each option is employed for its stability and computational advantages, so that each point on the surface is the volatility which zeroes the average $P\&L$ over 10 000 paths.

The volatility surface in Figure 3.1 is constructed for standard European options where the maturities range from three months to 18 months, increasing in increments of three months. The strike price ranges from 85% to 115% of the spot price $S_0 = 100$.

In practice this amount of data is not available, however when given enough data we see that calculating the break-even volatility can recover the true implied volatility surface fairly well. In Figure 3.1, the MSE between the implied and break-even volatility surfaces is 3.7873e-07. Heston (1993) acknowledges that daily returns are asymptotically normally distributed as the variance process reaches a steady state distribution in the long run with mean θ , given the mean reversion parameter $\kappa > 0$ (Cox *et al.*, 1985). This explains why the standard methodology works well, with Heston simulated data, for options with longer maturities and is able to overcome the instabilities frequently observed when estimating break-even volatility for short-term options. This is evident in Figure 3.1 where the MSE for the shortest skew is tenfold that of the MSE for the longest skew (1.2247e-06 and 1.8617e-07 respectively). This is a significant difference considering the sample size.




Fig. 3.1: True implied volatility (transparent) and break-even volatility estimated surface (10 000 paths, MSE: 3.7873e-07).

Monte Carlo option pricing typically relies on simulating a larger amount of data than used here, however the break-even method performs well regardless. If there is evidence that a certain stock follows a SV process or there is a degree of serial correlation in a share price, the amount of data used here would not be available. Even for the shorter term options where there is technically more data available due to shorter independent time series needed, instabilities arise and estimates worsen. A 20-year time series of daily prices would provide only 80 three-month non-overlapping paths, and thus other methods must be explored to improve the approximation.

3.3 Proposed method using Heston Greeks

In an attempt to improve the replication of the option in the delta-hedging process, we change the reference model from Black-Scholes to the correct model from which the data is being sourced. The aim is to obtain a more accurate picture of the volatility skew than the standard methodology and to assess whether the potential reduction in error from changing the reference model is worthwhile. Apart from an implied volatility price, another output is the Heston parameters which imply this price. Thus we can obtain the model dynamics as an added bonus, somewhat of a pseudo calibration.

Changing the hedging and pricing model substantially alters the scale of the problem. In the case of the Heston model, we now need to zero the $P\&L$ which

becomes a function of multiple parameters. This presents itself as an optimisation problem akin to market calibration of the Heston model, with accompanying difficulties. Furthermore, implementing a delta hedge using the Heston model requires the starting value of the variance process (ν_0) to be known. In our case of dynamic hedging along a stock price path, the corresponding latent variance process needs to be known at each point in order to calculate the Γ_{t_i} values in Equation (2).

The proposed way forward is to simulate data using the Milstein discretisation, thus allowing the volatility path to be known, which is not the case in practice. This would allow us to see the extent to which a change in the reference model reduces estimation error when compared to the standard BEV methodology.

With the $P\&L$ becoming a function of the parameter set $\{\kappa, \theta, \sigma, \rho\}$, instead of a single volatility σ with Black-Scholes delta hedging, zeroing the $P\&L$ becomes an optimisation problem. The initial variance ν_0 is usually a parameter of the Heston model, but since we are performing pure simulation and have access to the variance process, we treat this as known.

For a break-even estimate based on a single price path (S_t , with corresponding variance path v_t), the following procedure explains how the optimisation is carried out to zero the $P\&L$ calculated using the Heston model. The optimisation is scaled up over many historical periods similarly to the standard methodology by either averaging estimates for each period or zeroing the average $P\&L$. For a specific option $P\&L$, this optimisation is outlined below:

1. For a given parameter set $\{\kappa, \theta, \sigma, \rho\}$, price the option using the Heston model and convert this price into a Black-Scholes implied volatility σ which is used in Equation (2).
2. Calculate the $P\&L$ from Equation (2) using the Heston Γ_{t_i} 's which rely on the Heston parameter set above and corresponding ν_{t_i} values.
3. Solve for the Heston parameters which zeroes the $P\&L$ using a global optimisation algorithm.
4. Convert these optimal parameters into a Black-Scholes implied volatility, as in 1, to find the break-even volatility.

An initial guess $x_0 = \{\kappa_0, \theta_0, \sigma_0, \rho_0\}$ must be set when performing the optimisation. If an extensive global optimisation algorithm is used, the optimal parameters are fairly stable between different cycles of the optimisation. Different solutions may occur due to random scattering of test points from the initial guess involved in the global algorithm. Consequently, the optimisation avoids finding local minima solutions and an accurate initial guess is therefore not crucial. Alternatively if

a global algorithm is not used, the solution will depend on x_0 and will likely find a local minimum.

The break-even methodology requires the $P\&L$ to be zeroed, one cannot simply minimise the $P\&L$ function. The method used in this dissertation is to minimise the square of the $P\&L$ function which will still allow a root to be found, should one exist as a function of the four Heston parameters, and maintain differentiability.

The optimisation is carried out on steps 1 and 2 above. The option needs to be priced in each iteration of the optimisation to yield an implied volatility from the parameters which is used in Equation (2). This appears as σ^2 in the equation.

Utilising this method involves characteristic function pricing methods where integrals need to be approximated to price and find the Greeks under the Heston model. We use the Little Heston Trap formulation (see A.2.2) in our optimisation. This specification of the characteristic function is used to calculate the Γ for the Heston model (see A.2.3) as well as to find the Heston price implied volatility (σ^2 in Equation (2)) from the Heston parameters being optimised. Separate from the optimisation procedure, the Heston price is calculated again in step 4 using the optimal Heston parameters and this price is converted into the final break-even volatility implied by the Heston model.

The variance process is simulated alongside the stock price process, rendering this investigation as a purely theoretical study. We simulate data using the Milstein discretisation with the same parameters as in Section 3.2, however with far fewer sample paths, and compare the standard method to the method described above. Following this, using the same data, we use the realised variance $(\frac{\Delta S_t}{S_t})^2$ as a proxy in place of the actual variance process ν_t and assess its viability as a substitute. This eliminates the need to have calibrated daily to uncover ν_t corresponding to a stock (given there is sufficient stability of the calibrated parameters), which can be used in the method above.

This method requires past price data only and may prove useful when no option information is available on the specific stock in order to calibrate the model to prices. However, if the model has been calibrated over time, then there will be a time series of daily parameters which can be used to obtain estimates of implied prices for options. These parameters only minimise the model pricing error to past observed prices and there is no implication that these would give true fair prices, which break-even volatility aims to do.

Chapter 4

Implementation and Results

In this section we construct BEV surfaces, initially using the Black-Scholes model to hedge, and then implement the methodology using the Heston model as the new reference model. The data is also sourced from the Heston model and the volatility surfaces are assessed by their MSE's. The intention is to see whether the change in reference model from the standard methodology improves estimation of the true implied volatility surface.

This is done with simulating Heston stock price paths with the same parameters as in Section 3.2. Both methods are computationally expensive, however the optimisation required when hedging with the Heston model slows the break-even method down even further and the usage of large datasets becomes infeasible. Therefore, only 50 independent paths are used for each method. We also only consider one-year options in order to limit the data size needed. Using fewer paths links to the fact that practitioners would not be able to bootstrap the data and create more paths, since this would break the covariance structure of the time series. Using 50 independent three-month stock price paths equates to a 12.5 year unbroken time series, this appears to be in the middle ground for amount of historical data. Any longer than is, there is a risk of static surfaces not changing over time and any shorter would not allow enough data. Naturally, for the one-year option, 50 independent one-year paths would not be readily available in most cases. To overcome this and have more sample paths, a practitioner may use paths which overlap. In doing this, it maintains the completeness of a path. However, if one is trying to obtain a distribution of break-even volatilities with the standard methodology (as in Section 2.1 and 2.3.3) then these volatility estimates for each path will no longer be *i.i.d.*. Applying overlapping time series means that a number of share price returns from one time series which form part of the weighted average of one break-even volatility estimate will be used in another estimate. The effect of which is not investigated in this dissertation and may not make a significant difference depending on the degree of overlapping. We choose to focus on using independent

time series for each break-even estimate.

4.1 Standard BEV

Here we see how the standard method performs on few sample paths. If we think of break-even volatility in the context of Monte Carlo, the more sample paths used, the closer the result will be to the true solution. We saw this in our Monte Carlo experiment in Figure 3.1 that, given enough sample paths, break-even volatility can recover the true implied volatility surface.




Fig. 4.1: Optimising the whole surface: True implied volatility surface (translucent) and break-even volatility surface, averaged over all options (50 paths). MSE = 3.0074e-05.

The result of using only 50 paths can be seen in Figure 4.1, where the average $P\&L$ for all options on the surface is zeroed. The common method is to calculate and zero the average $P\&L$ for one option over many time series. Here we have modified this approach and averaged all options' $P\&L$ s over all of the time series. In tests between finding the whole surface at once versus for each option separately, there were significant decreases in computational time using the former method. However, averaging all options' $P\&L$ s and zeroing this average does not necessarily imply that each option's average $P\&L$ is zeroed. In fact, not having this requirement yields an inaccurate representation of the volatility surface as different option's $P\&L$ s can offset each other in this average and result in wildly inaccurate break-even volatility estimates. In our implementation, the MATLAB minimiser `fmincon` is used where the input variable is the whole volatility surface/matrix.

Each element of the matrix represented a break-even volatility for the corresponding option on the surface. The volatilities are used to calculate each option's average $P\&L$ over the 50 time series and then the whole surface $P\&L$ is averaged. A non-linear condition was added onto the optimisation to ensure each option's average $P\&L$ was zeroed and the minimisation was carried out on the square of the average $P\&L$ of the whole surface to ensure differentiability in finding the root of this function. In turn, this leads to the same solution as if you were to construct the volatility surface with each option individually by zeroing the average $P\&L$, although with advantageous computing time.

For comparison we've added Figure 4.2 where the break-even volatility is found for each path and each option, producing a distribution of volatilities. The surface is well behaved and lies within the three standard deviation bounds. In most cases we find that zeroing the average $P\&L$ produces smoother results and improves convergence to the true result. However, the MSE of 4.2 is similar in scale to that of Figure 4.1, with the added benefit of having a certain degree of confidence in the estimate due to the resultant distribution.




Fig. 4.2: Point-wise optimisation: True implied volatility surface (transparent) and break-even volatility surface, zeroing each path's $P\&L$ for a specific option separately (50 paths), in between 3 standard deviation bounds (red). MSE = $2.2235e-05$.

In Figure 4.3, the break-even volatility distributions for two options (one-year term, strikes of 85 and 100) have been included as an example of what one would obtain by calculating the break-even volatility for each sample path. The final BEV estimates are either the means of the distributions or the volatilities which zero

the average $P\&L$'s, the latter method being more accurate as seen in Table 4.1. With such few sample paths, the estimates in Table 4.1 are slightly lower than the true implied volatility as expected. The distributions suffer from significant sample variation due to the small sample size, this is evident from the instabilities and multiple peaks. However, we do obtain accurate confidence intervals (Figure 4.2) with the distributions being seemingly symmetrical. This makes the method of finding the distribution more attractive, as the highly volatile paths will be quantified as an outlying volatility in the distribution.




Fig. 4.3: Break-even volatility distributions for two options of 1-year term (sample size: 50).

	$(K = 85, \tau = 1)$	$(K = 100, \tau = 1)$
True implied volatility	0.2517	0.2447
Zero average $P\&L$	0.2461	0.2424
Distribution mean	0.2444	0.2422

Tab. 4.1: Break-even volatility estimates for two options (50 sample paths).

4.2 Optimising over Heston parameters

In initial tests, point-wise estimation of break-even volatilities is performed. This means that the surface is constructed option-by-option and the Heston parameters, which imply the break-even volatility, were found by zeroing the average $P\&L$ of all sample paths. Since Γ and σ^2 in Equation (2) are now functions of the parameter

set $\{\kappa, \theta, \sigma, \rho\}$, we need to calculate these by approximating the integrals in characteristic function pricing. This is done by mid-point quadrature shown in Appendix A.2.2 to calculate the Heston price implied volatility σ and the Heston Γ in A.2.3. To implement the break-even method using Heston parameters, one needs to find the corresponding Γ values along a price and volatility path, this involves a summation to approximate the integral at each point along a time series. Following this, the sum needs to be taken over each path to find the associated $P\&L$ and then averaged over all the paths. The most effective implementation of this involves some nifty matrix manipulation and multiplication in MATLAB which will substantially speed up computational time and be useful when optimising over the full implied volatility surface. In our case, we simulate data where each stock path forms a row of the data matrix. We then expand into three dimensions to perform the quadrature at each point along a path.

The optimisation presents difficulties in that it is now similar to the problem of market calibration of the Heston model where an objective function needs to be minimised over the parameters as well. In calibration of the Heston model, different parameter sets can yield the same value of the objective function giving evidence that it, and the $P\&L$ function under consideration, is flat at the minimum (Cui *et al.*, 2017). In our case, different combinations of the Heston parameters can be returned as solutions which minimise the $P\&L$. The optimal solution may depend on the optimisation algorithm and several sets of solutions may be found due to different starting points leading to local minima. Fortunately, we have found that various parameter sets all which minimise the $P\&L$ close to zero yield the same implied volatility around the correct level. This method is able to pick up the shape of the skew, however the optimal parameters may not be the true parameters which simulated the data.

The parameter set $\{\kappa, \theta, \sigma, \rho\}$ solution is sensitive to the starting point chosen. To circumvent long run times, the function `fmincon` is used with the starting point as the true solution $\{\kappa, \theta, \sigma, \rho\} = \{9, 0.06, 0.5, -0.4\}$. In cases where no information is available to choose an appropriate initial point, a global optimisation algorithm (e.g. `gs` in MATLAB) can be used with similar effect.

The result of this point-wise optimisation can be seen in Figure 4.4 where the MSE is similar in scale to those of the standard method in Figures 4.1 and 4.2. We know that implementing the theoretically correct hedge, by using the Heston Greeks, should yield a more accurate implied volatility solution and converge quicker with more data. With the small sample size, we can only note that a change in the reference model works just as well.

The $P\&L$ function becomes extremely flat at the minimum or rather when the




Fig. 4.4: Point-wise optimisation: True implied volatility surface (transparent) and point-wise optimised for Heston parameters, zeroing each options' average $P\&L$ (50 paths). $MSE = 2.6976e-05$.

squared $P\&L$ is close to zero, since this is the function being minimised. In Figure 4.5, the MSE is calculated for the 3-month skew with increasing sample size, where the data is simulated with the same parameters as before. Taking the implementation a step further, the whole skew is optimised to give one parameter set, not as before by calculating the $P\&L$ option by option. After some large variations,




Fig. 4.5: MSE as a function of sample size for the 3-month volatility skew, with strike price range 85% – 115% of S_0 .

the MSE appears to decrease and flatten out with no further improvement and the parameter estimates in Table 4.2 paints the same picture.

Parameters	Sample size (number of paths)					True values
	100	200	300	400	500	
κ	8.9885	8.9938	8.9934	8.9941	8.9930	9
θ	0.0552	0.0596	0.0591	0.0589	0.0606	0.06
σ	0.5185	0.5173	0.5195	0.5157	0.5186	0.5
ρ	-0.5111	-0.5042	-0.5129	-0.5039	-0.5114	-0.4

Tab. 4.2: Parameter estimates for 3-month skew with increasing sample size.

It appears that the estimates for κ , θ and σ are close to the true values, but admittedly the success of this estimation highly depends on the choice of parameters used to simulate and the interval bounds used in the minimisation algorithm of the $P\&L$. Restricting the bounds for ρ to be between $[-1, 0]$, which fits for most equity option prices, allows for more accurate estimates as seen in Table 4.2, otherwise with $[-1, 1]$ bounds for ρ the estimates are even further astray. The parameter θ can be estimated with the most confidence, regardless of the interval bounds or the chosen true value used in the simulation. This is the least we could hope for as the implied volatilities can be somewhat approximated using the Heston model as seen in Figure 4.4. The long-term variance θ controls the vertical position of the volatility surface and since this method can recover these average volatility levels as well as the approximate shape of the skew, it makes sense that θ can be reliably estimated.

Analogous to market calibration which uses option of multiple strikes and maturities, attempting to construct the whole break-even volatility surface at once is superior. Not only does this constrain the optimisation problem further, it also creates a smoother surface as the output will be one set of parameters i.e. one model, not one for each option or skew. Calculating the $P\&L$ of the entire surface, in conjunction with the characteristic function pricing, requires one to use higher data dimensions (4th and 5th in MATLAB) to avoid loops. The MSE of the surface in Figure 4.6 is significantly better than optimising point-wise with the Heston model or the standard methodology using Black-Scholes. The parameters which minimise the surface's $P\&L$ over all simulated paths are $\{\kappa, \theta, \sigma, \rho\} = \{7.9079, 0.0601, 0.5121, -0.5054\}$. These parameter estimates are in line with previous estimates based on single options or skews based on one maturity, with the exception of κ . This method does become arduous on both the user side and computationally with the optimisation, however, the estimation performs rather

remarkably condising the small sample size. If outright accuracy is required, this method provides a smaller MSE than the standard break-even methodology and we expect it to decrease with a larger dataset.




Fig. 4.6: Optimising the whole surface: True implied volatility surface (transparent) and full surface optimised for Heston parameters, averaging all options' $P\&L$'s (50 paths). $MSE = 8.6556e-06$.

4.3 Incorporating realised volatility

Changing the reference model should in theory yield an improved result, however in practice the actual variance process ν_t is not observable. Due to this the methods outlined thus far using the Heston model cannot be used in their entirety.

In the Heston model we have the following relationship:

$$\left(\frac{dS_t}{S_t}\right)^2 = \nu_t dt$$

This is true in a continuous-time setting, however we test this in a discrete setting and substitute for the realised variance where the actual variance process ν_t has been needed in the previous section.

In doing this, the initial variance ν_0 is added to the problem as another variable which needs to be optimised. The increased dimensionality of the optimisation and using the realised variance instead severely worsens estimates of parameters and thus the final break-even surface. The amount of price information without the true variance path is insufficient to obtain accurate parameter estimates. This is

alleviated if one is at least able to estimate the initial variance ν_0 to reduce the scale of the optimisation.

The most accurate results obtained are in Figure 4.7 using point-wise estimation of the surface and optimised over all parameters, including ν_0 . As anticipated, the MSE worsens yet the surface is similar to that in Figure 4.4 and arguably does surprisingly well using this crude proxy for the variance process.

Further work can be done along the same lines as calibration of the Heston model where the employment of dimension reduction techniques can be used to simplify the optimisation. Prior knowledge can be used to estimate certain parameters, eliminating them from the optimisation. As an example, amongst other heuristics discussed by Cui *et al.* (2017), this can include ν_0 which is usually set to the short-term ATM implied variance. This may greatly help where the variance path is not available in the form of daily calibrated values for ν_t , which could be used if the calibrated parameters are fairly constant, or where the realised variance is used as a proxy.




Fig. 4.7: True implied volatility surface (transparent) and point-wise optimised over Heston parameters (50 paths) using the realised variance in place of the actual variance path. MSE = 1.0397e-04.

Chapter 5

Discussion and Conclusion

The main question posed in this dissertation is whether it is beneficial to change the reference model in the break-even volatility calculation from Black-Scholes to another. The chosen stochastic volatility model under examination is troublesome to work with in the BEV methodology, due to the pricing methods available as well as the method's inability to use large amounts of data.

The recommendation when faced with limited data, where one supposes there may be some serial correlation amongst prices, would be to calculate the break-even volatilities for each historical period in order to produce a distribution of volatility estimates. The existence of outliers may help update pricing decisions and help load in any volatility risk premium.

The standard Black-Scholes methodology tends to underprice most options as seen in Figures 4.1 and 4.2. This is fairly prevalent with little data. The fact that this occurred in Figure 3.1, where the break-even surface approached the true implied volatility surface mostly from beneath, means this might be a systematic issue for the standard methodology. Changing the hedging model to Heston appears to make the spread between underpricings and overpricings more level, this is seen in the figures in Section 4.2.

Apart from an accurate implied price, changing the reference model to Heston can allow a practitioner to better understand the risks being faced even when pricing and hedging an option. Mostly, the shape of the skew, given by σ and ρ , can give an idea of these risks. In the implementation in this dissertation, it was found that κ and θ were the most reliable to estimate. This means a practitioner could still acquire an idea of the average volatility level and how long the volatility may stay above or below this level.

As a closing remark, using a different reference model in the break-even volatility method means the $P\&L$ will not be of the same form as Equation (1). There is still residual risk that is not hedged even by changing the reference model. However, the normal Black-Scholes method cannot account for this extra volatility risk,

arising from Heston data, in the same way a simple change in reference model cannot. Using the correct hedge ratios can at least offset some risk from the additional volatility in the stock price. We have seen that using the Heston model with Heston simulated data replicates the option $P\&L$ and calculates break-even volatilities comparably to the standard method, by implementing the correct hedge and provides a pseudo market calibration.

Bibliography

- Albrecher, H., Mayer, P., Schoutens, W. and Tistaert, J. (2007). The Little Heston Trap, *Wilmott* (1): 83–92.
- Bakshi, G., Cao, C. and Chen, Z. (1997). Empirical Performance of Alternative Option Pricing Models, *The Journal of finance* **52**(5): 2003–2049.
- Black, F. and Scholes, M. (1973). The Pricing of Options and Corporate Liabilities, *Journal of Political Economy* **81**(3): 637–659.
- Cox, J. C., Ingersoll, J. E. and Ross, S. A. (1985). A Theory of the Term Structure of Interest Rates, *Econometrica* **53**: 385–407.
- Cui, Y., del Baño Rollin, S. and Germano, G. (2017). Full and fast calibration of the Heston stochastic volatility model, *European Journal of Operational Research* **263**(2): 625–638.
- Dupire, B. (2006). Fair Skew: Break-Even Volatility Surface, *Technical report*, Bloomberg L.P.
- Gatheral, J. (2011). *The Volatility Surface: A Practitioner's Guide*, Vol. 357, John Wiley & Sons.
- Gil-Pelaez, J. (1951). Note on the Inversion Theorem, *Biometrika* **38**(3-4): 481–482.
- Heston, S. L. (1993). A Closed-Form Solution for Options with Stochastic Volatility with Applications to Bond and Currency Options, *The Review of Financial Studies* **6**(2): 327–343.
- Rouah, F. D. (2013). *The Heston Model and Its Extensions in Matlab and C#*, John Wiley & Sons.
- Rouah, F. D. and Vainberg, G. (2007). *Option Pricing Models and Volatility Using Excel-VBA*, Vol. 361, John Wiley & Sons.
- Suzuki, K. and Vyas, K. (2011). Break-even volatility: A hedging cost approach for option volatility, *Technical report*, Nomura Securities.

Appendix A

Model specifics

A.1 Black-Scholes PDE

The Black-Scholes partial differential equation as it appears in [Black and Scholes \(1973\)](#) for a European option C with risk-free rate r and volatility σ .

$$\frac{\partial C}{\partial t} + rS \frac{\partial C}{\partial S} + \frac{1}{2} \sigma^2 S^2 \frac{\partial^2 C}{\partial S^2} - rC = 0 \quad (\text{A.1})$$

In terms of the Greeks, the equation for Θ_t as substituted in Section [2.2.1](#) becomes

$$\Theta_t = rC - rS\Delta_t - \frac{1}{2} \sigma^2 S^2 \Gamma_t \quad (\text{A.2})$$

A.2 The Heston stochastic volatility model

A.2.1 Simulation

There are a number of discretisation schemes at one's disposal to simulate the Heston stock and volatility path. We will be using the Milstein scheme to simulate data which produces fewer negative values for the variance process than the Euler scheme ([Rouah, 2013](#), p. 183–185), and a truncation scheme in case the simulated variance is negative.

The stock price path is given by

$$S_i = \begin{cases} S_0 & \text{if } i = 0 \\ S_{i-1} \exp \left((\mu - \frac{1}{2} \nu_{i-1}) \Delta t + \sqrt{\nu_{i-1}} \Delta t Z_{1,i} \right) & \text{if } i > 0 \end{cases}$$

where

$$\nu_i = \begin{cases} \nu_0 & \text{if } i = 0 \\ \left(\nu_{i-1} + \kappa(\theta - \nu_{i-1}) \Delta t + \sigma \sqrt{\nu_{i-1}} \Delta t Z_{2,i} + \frac{1}{4} \sigma^2 (Z_{2,i}^2 - 1) \Delta t \right)^+ & \text{if } i > 0 \end{cases}$$

$Z_{1,i}$ and $Z_{2,i}$ are $N(0, 1)$ random variables generated with correlation ρ .

In most sources the schemes are given in risk-neutral form, however we need to simulate real-world data as used by the break-even methodology i.e. with drift parameter μ .

A.2.2 Characteristic function pricing

In this dissertation, we have been primarily concerned with pricing and finding the break-even volatilities for European call options. The price of a European call C with strike K and maturity T on a stock S can be written as

$$C = S_0 \mathbb{E}^{\mathbb{Q}_S} [\mathbb{I}_{\{S_T > K\}}] - \frac{K}{A_T} \mathbb{E}^{\mathbb{Q}} [\mathbb{I}_{\{S_T > K\}}]$$

where $A_t = e^{rt}$ is the numeraire/bank account, \mathbb{Q} is the risk-neutral measure and \mathbb{Q}_S is the measure given by

$$\frac{d\mathbb{Q}_S}{d\mathbb{Q}} = \frac{S_T/S_0}{A_T/A_0}.$$

The expectations of the indicator functions are just measures of probability of the event $\{S_T > K\}$ which we call P_1 and P_2 , thus

$$\begin{aligned} C &= S_0 \mathbb{Q}_S(\ln(S_T) > \ln(K)) - K e^{-rT} \mathbb{Q}(\ln(S_T) > \ln(K)) \\ &= S_0 P_1 - K e^{-rT} P_2. \end{aligned}$$

To find P_1 and P_2 we use a special case of the inversion formula generalised by [Gil-Pelaez \(1951\)](#), which says for $k \in \mathbb{R}$ and random variable X with associated characteristic function ϕ_X , then

$$Pr(X > k) = \frac{1}{2} + \frac{1}{\pi} \int_0^\infty \Re \left[\frac{e^{-iuk} \phi_X(u)}{iu} \right] du$$

where \Re denominates the real part of the integrand and $\phi_X(u) = \mathbb{E}[e^{iuX}]$ is the characteristic function.

Working under the risk-neutral measure to find the characteristic function for the terminal log-stock price $s_T = \ln(S_T)$, letting $k = \ln(K)$ and approximating the integrals using simple quadrature (midpoint), this gives

$$\begin{aligned} P_1 &= \frac{1}{2} + \frac{1}{\pi} \sum_{n=1}^N \Re \left[\frac{e^{-iu_n k} \phi_{s_T}(u_n - i)}{iu_n \phi_{s_T}(-i)} \right] \delta u \\ P_2 &= \frac{1}{2} + \frac{1}{\pi} \sum_{n=1}^N \Re \left[\frac{e^{-iu_n k} \phi_{s_T}(u_n)}{iu_n} \right] \delta u \end{aligned}$$

where the integration is limited to the range $[0, u_{\max}]$, $\delta u = u_{\max}/N$ and $u_n = (n - \frac{1}{2})\delta u$.

Note: The integrands for P_1 and P_2 must be inspected to choose a suitable u_{\max} and number of quadrature points N .

The Little Heston Trap formulation

In the introductory paper on the model, [Heston \(1993\)](#) provided a characteristic function pricing approach, however a method specified by [Albrecher et al. \(2007\)](#) which proved to be more numerically stable was used in this dissertation.

The characteristic function of the log-stock price $s_T = \ln(S_T)$ under \mathbb{Q} is

$$\phi_{s_T}(u) = \exp(C + D\nu_0 + iu\ln(S_0))$$

where

$$C = rTiu + \theta\kappa \left(Tx_- - \frac{1}{a} \ln \left(\frac{1 - ge^{-Td}}{1 - g} \right) \right)$$

$$D = \frac{1 - e^{-Td}}{1 - ge^{-Td}} x_-$$

and

$$a = \frac{1}{2}\sigma^2 \quad b = \kappa - \rho\sigma iu \quad c = -\frac{u^2 + iu}{2}$$

$$d = \sqrt{b^2 - 4ac} \quad x_{\pm} = \frac{b \pm d}{2a} \quad g = \frac{x_-}{x_+}$$

A.2.3 The Heston Greeks

From above, we have the price of a call at time t as

$$C_t = S_t P_1 - K e^{-r(T-t)} P_2$$

As in the Black-Scholes case, we take our delta as

$$\Delta_t = P_1$$

This is shown in [Rouah and Vainberg \(2007, p. 203\)](#) and also used by [Bakshi et al. \(1997\)](#).

The gamma is found by taking the derivative again. In the Little Trap formulation, working with the summation for P_1 , this is

$$\Gamma_t = \frac{\partial P_1}{\partial S_t}$$

$$= \frac{1}{\pi} \sum_{n=1}^N \Re \left[\frac{e^{-iu_n k}}{iu_n} \frac{\partial}{\partial S_t} \left[\frac{\phi_{s_T}(u_n - i)}{\phi_{s_T}(-i)} \right] \right] \delta u$$

By the product and chain rule, the derivative inside the summation becomes

$$\frac{\partial}{\partial S_t} \left[\frac{\phi_{s_T}(u_n - i)}{\phi_{s_T}(-i)} \right] = \frac{\phi_{s_T}(u_n - i) \frac{i(u_n - i)}{S_t} \phi_{s_T}(-i) - \phi_{s_T}(u_n - i) \frac{\phi_{s_T}(-i)}{S_t}}{[\phi_{s_T}(-i)]^2}$$

$$= \frac{iu_n}{S_t} \frac{\phi_{s_T}(u_n - i)}{\phi_{s_T}(-i)}$$

Therefore,

$$\Gamma_t = \frac{1}{\pi S_t} \sum_{n=1}^N \Re \left[e^{-iu_n k} \frac{\phi_{s_T}(u_n - i)}{\phi_{s_T}(-i)} \right] \delta u$$